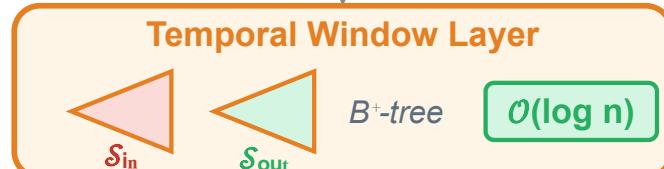
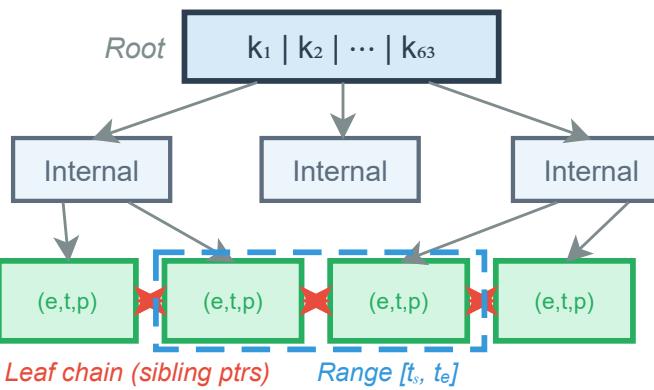


(a) Three-Layer Index Architecture



Robin Hood hashing for $O(1)$ expected lookup

(b) B⁺-Tree Structure



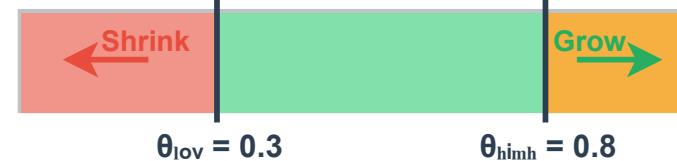
Entry: (edge_id, timestamp, attr_ptr) \rightarrow 24 bytes

Thm 1: Insert $\in \mathcal{O}(\log n)$ Thm 2: Query $\in \mathcal{O}(\log n + k)$

Empirical $R^2 > 0.97$

$B = 64$

(c) Adaptive Sequence Management

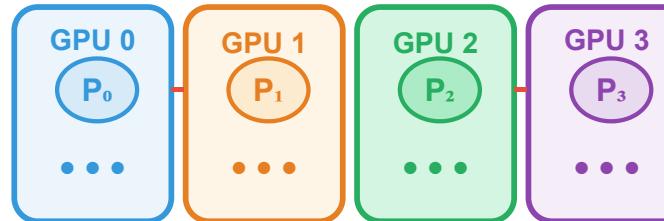


Sequence Bounds
 $L_{\min} = 8 \leq |\mathcal{S}(v)| \leq L_{\max} = 512$
 Growth factor: $\alpha = 0.2$

Performance Gains
Memory Reduction: 23.7%
F1 Improvement: +0.41%

$U < \theta_{lov}$ Optimal $U > \theta_{hi_mh}$

(d) Locality-Aware Partitioning



Cross-partition edges: 8.3% (vs 34.7% hash)

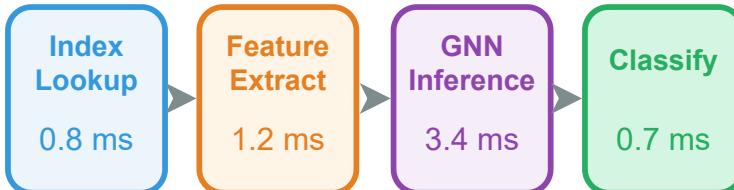
$$|P_i| \in [(1-\beta)\bar{P}, (1+\beta)\bar{P}], \beta = 0.15$$

Hash-based: 34.7%

Ours: 8.3% ↓ 76% reduction

$$\text{Similarity: } \sigma(v, P_i) = |\mathcal{N}(v) \cap P_i| / |\mathcal{N}(v)|$$

(e) End-to-End Latency Breakdown



P50 6.1 ms Target: < 50 ms ✓ **P99** 18.7 ms

P50 P95 P99

Latency distribution (median \rightarrow 99th percentile)

(f) Fault Tolerance & Recovery

Two-Phase Commit Checkpointing



Every 500 iterations (~ 300 s)

Recovery Phases (Total: 40.9 s)

Model 8.4s	Part. 6.3s	Replay (TAIL): 22.5s
------------	------------	----------------------

100% Recovery Success
 Mean: 40.9 s | Target: < 60 s ✓

Delta encoding: 73% size reduction | 4.7 GB/s bandwidth